

E N T 3 5 6 - S

# How Cox Automotive Runs GitHub Enterprise on AWS

Bryan Cross  
Principal Solutions Engineer  
**GitHub**

@bryancross

Felipe Ramirez  
Senior Systems Engineer  
**Cox Automotive**

@felipe-ramirez



Bryan Cross  
Principal Solutions Engineer  
GitHub







**1,100,000,000**

Contributions this year

**31,000,000**

Developers

100,000,000

Repositories

8,000,000

New developers

# 80%

Of users are from outside U.S.

# What's a Git?

# What's a Git?



# What's a Git?

git

/git/

*noun* BRITISH *informal*

an unpleasant or contemptible person.

# What's a Git?

“If you can’t figure out who the Git on your team is, it’s probably you”

- *Your Team*





\$41,500,000,000



\$41,500,000,000



# Built for developers

GitHub is a development platform inspired by the way you work. From **open source** to **business**, you can host and review code, manage projects, and build software alongside 31 million developers.

Username

 Pick a username

Email

 you@example.com

Password

 Create a password

Make sure it's more than 15 characters, or at least 8 characters, including a number, and a lowercase letter.

[Sign up for GitHub](#)

By clicking "Sign up for GitHub", you agree to our [terms of service](#) and [privacy statement](#). We'll occasionally send you account related emails.

Same UX as GitHub

Scalable

Easy to  
manage/upgrade

Built in HA/DR

A true platform





Bryan Cross  
Principal Solutions Engineer  
GitHub





Bryan Cross  
Principal Solutions Engineer  
GitHub





Bryan Cross  
Principal Solutions Engineer  
GitHub

SALES GUY



Events

Tags

Reports

Limits

**INSTANCES**

Instances

Launch Templa

Spot Requests

Reserved Instan

Dedicated Host

Scheduled Insta

Capacity Reser

**IMAGES**

AMIs

Bundle Tasks

**ELASTIC BLOCK**

Volumes

Snapshots

Lifecycle Mana

**NETWORK & SEC**

Security Groups

Elastic IPs

Placement Gro

## Step 1: Choose an Amazon Machine Image (AMI)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance.

You can select an AMI provided by AWS, our user community, or the AWS Marketplace; or you can select one of your own AMIs.

Search for an AMI by entering a search term e.g. "Windows"



### Quick Start (0)

My AMIs (2)

AWS Marketplace (1187)

Community AMIs (6390)

Free tier only (i)



[Cancel and Exit](#)

**THIS BAD BOY CAN HOLD**

**GitHub Enterprise**

**SO MANY REPOS**





Felipe Ramirez  
Senior Systems Engineer  
Cox Automotive



INVENTORY  
SOLUTIONS

Manheim

RMS/AUTOMOTIVE



汽车街  
[autostreets.com](http://autostreets.com)



**carsguide**



[Dealer-Auction.com](http://Dealer-Auction.com)



**DEALER SOLUTIONS**

**精真估**  
[jingzhengu.com](http://jingzhengu.com)



RETAIL  
SOLUTIONS

HOMENET  
AUTOMOTIVE

Dealertrack 

VinSolutions

DEALER.COM

oxtime

vAuto  
LIVE MARKET VIEW



Kelley Blue Book



Autotrader



FINANCIAL  
SOLUTIONS

**NEXTGEAR**  
CAPITAL™

Mahindra  
**First Choice**

Modix 

**MOLICAR**

**MOTORS**  
.co.uk

**MOVE**X



MOBILITY  
SOLUTIONS

Getaround

RIDECELL



CLUTCH



FLEXDRIVE



OUSTER



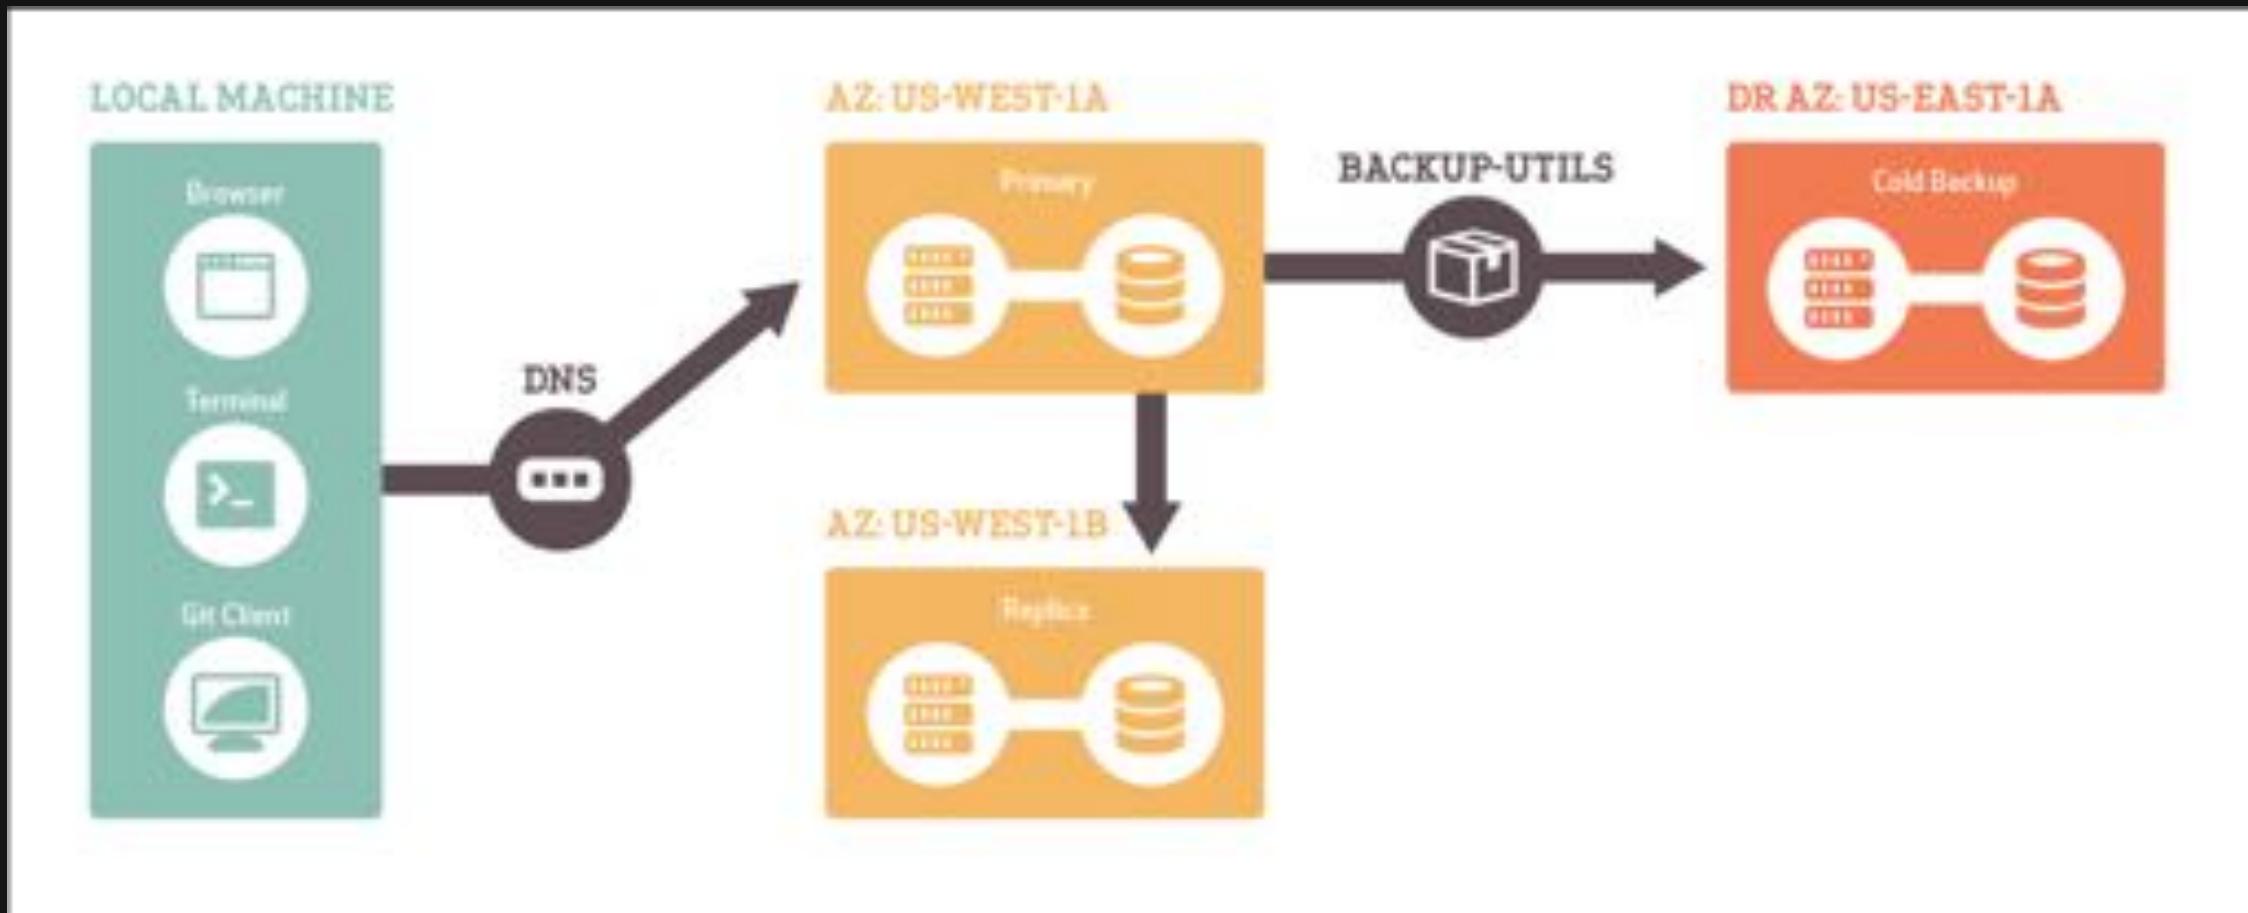
[github.com/felipe-ramirez/reinvent18](https://github.com/felipe-ramirez/reinvent18)

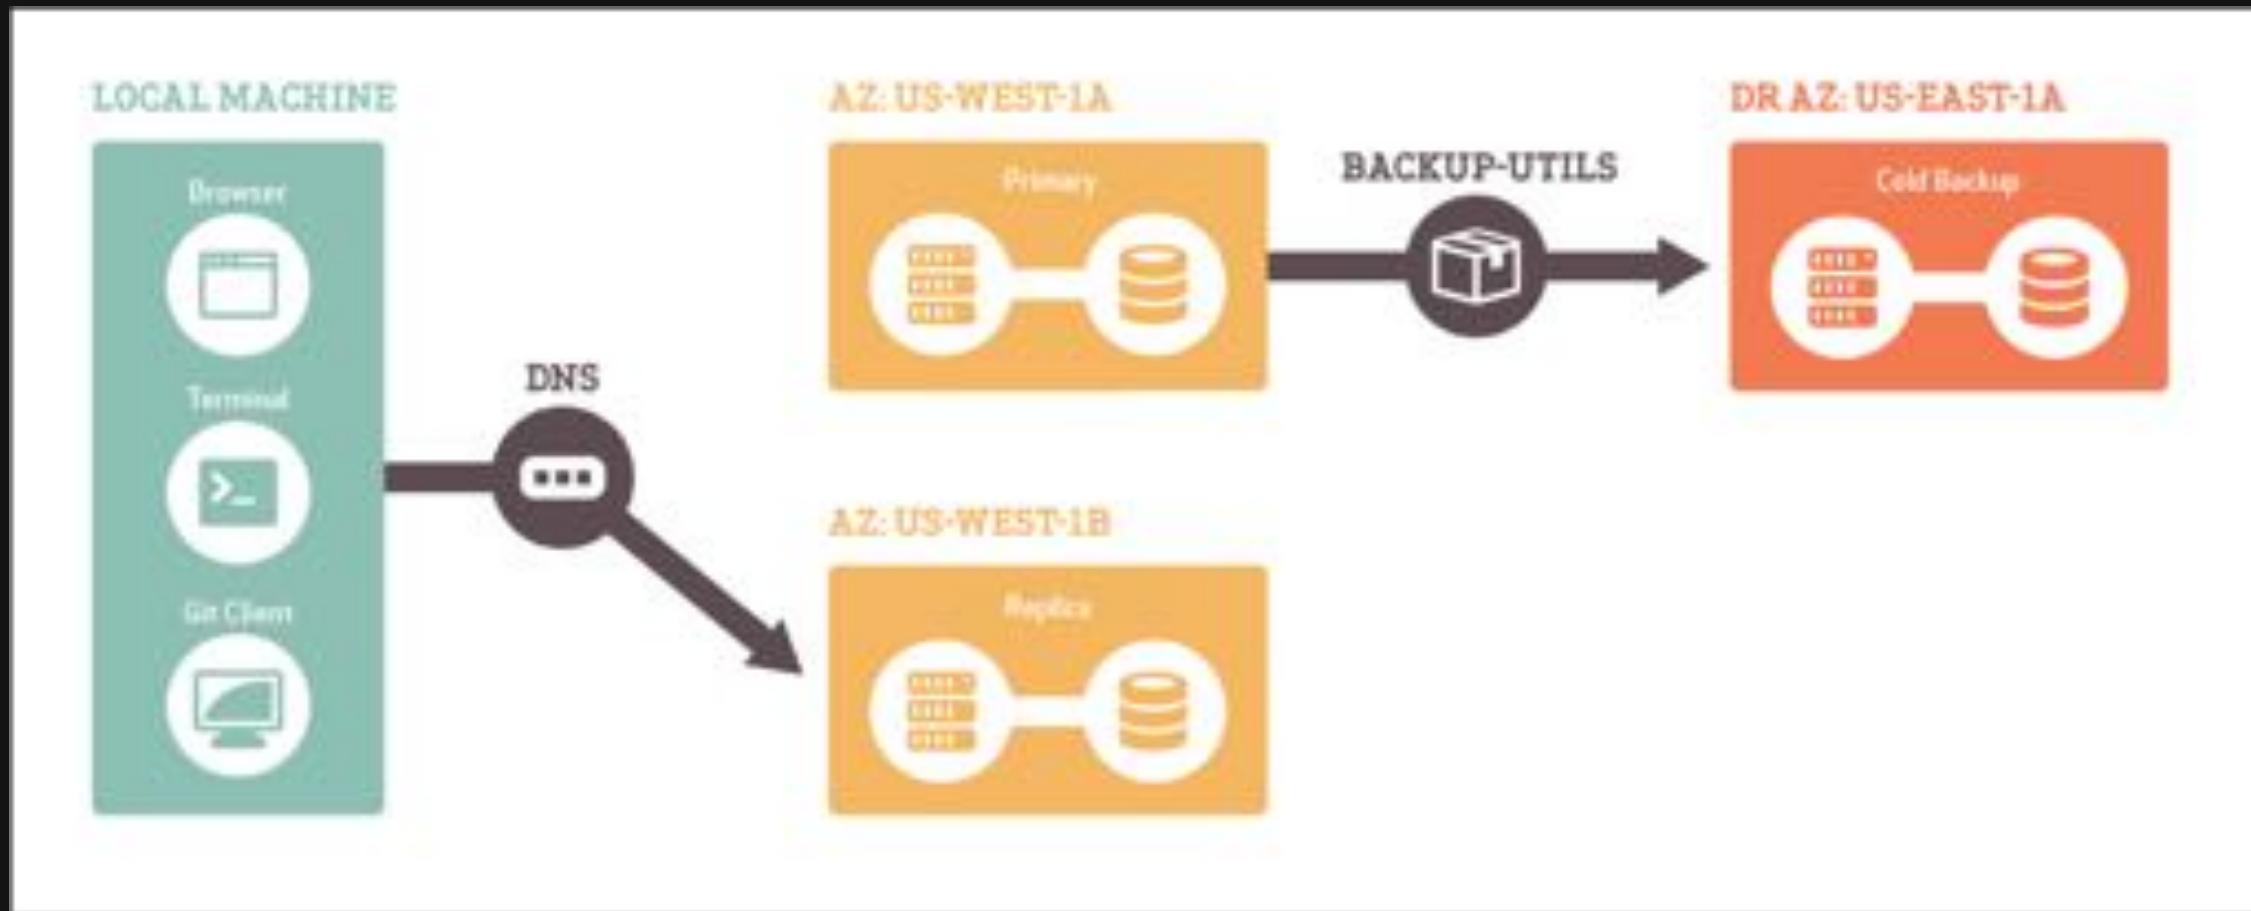
# Agenda

1. GitHub Enterprise HA on AWS
2. Defining infrastructure with Terraform
3. Using API & Terraform for configuration
4. Future enhancements

# Agenda

1. GitHub Enterprise HA on AWS
2. Defining infrastructure with Terraform
3. Using API & Terraform for configuration
4. Future enhancements





HA doesn't  
replace backups



[github.com/github/backup-utils](https://github.com/github/backup-utils)

# Agenda

1. GitHub Enterprise HA on AWS
2. Defining infrastructure with Terraform
3. Using API & Terraform for configuration
4. Future enhancements

Why deploy GitHub  
Enterprise with code?

Code as  
documentation

Consistency  
and speed

It's fun to do



HashiCorp

Terraform



# GitHub Enterprise AMI ID (ami.tf)

```
data "aws_ami" "ghe" {
  filter {
    name    = "name"
    values = ["GitHub Enterprise ${var.ghe_version}"]
  }

  owners = ["895557238572"]
}
```

# GitHub Enterprise AMI ID (ami.tf)

```
data "aws_ami" "ghe" {
  filter {
    name    = "name"
    values = ["GitHub Enterprise ${var.ghe_version}"]
  }

  owners = ["895557238572"]
}
```

# GitHub Enterprise AMI ID (ami.tf)

```
data "aws_ami" "ghe" {
  filter {
    name    = "name"
    values = ["GitHub Enterprise ${var.ghe_version}"]
  }
  owners = ["895557238572"]
}
```

# EC2 instances (main.tf)

```
resource "aws_instance" "ghe" {
    count          = 2
    ami            = "${data.aws_ami.ghe.id}"
    subnet_id      = "${element(var.subnet_ids, count.index)}"
    instance_type  = "${var.instance_type}"
    vpc_security_group_ids = ["${aws_security_group.ghe.id}"]

    ebs_block_device {
        device_name = "/dev/xvdd"
        volume_type = "gp2"
        volume_size = "${var.data_volume_size}"
    }

    tags = {
        Name = "${local.env_tag}-${format("%01d", count.index + 0)}"
    }
}
```

# EC2 instances (main.tf)

```
resource "aws_instance" "ghe" {
    count          = 2
    ami           → = "${data.aws_ami.ghe.id}"
    subnet_id      = "${element(var.subnet_ids, count.index)}"
    instance_type   = "${var.instance_type}"
    vpc_security_group_ids = ["${aws_security_group.ghe.id}"]

    ebs_block_device {
        device_name = "/dev/xvdd"
        volume_type = "gp2"
        volume_size = "${var.data_volume_size}"
    }

    tags = {
        Name = "${local.env_tag}-${format("%01d", count.index + 0)}"
    }
}
```

# EC2 instances (main.tf)

```
resource "aws_instance" "ghe" {
    count          = 2
    ami            = "${data.aws_ami.ghe.id}"
    subnet_id      = "${element(var.subnet_ids, count.index)}"
    instance_type  = "${var.instance_type}"
    vpc_security_group_ids = ["${aws_security_group.ghe.id}"]

    ebs_block_device {
        device_name = "/dev/xvdd"
        volume_type = "gp2"
        volume_size = "${var.data_volume_size}"
    }

    tags = {
        Name = "${local.env_tag}-${format("%01d", count.index + 0)}"
    }
}
```

# EC2 instances (main.tf)

```
resource "aws_instance" "ghe" {
    count          = 2
    ami            = "${data.aws_ami.ghe.id}"
    subnet_id      = "${element(var.subnet_ids, count.index)}"
    instance_type  = "${var.instance_type}"
    vpc_security_group_ids = ["${aws_security_group.ghe.id}"]

    ebs_block_device {
        device_name = "/dev/xvdd"
        volume_type = "gp2"
        volume_size = "${var.data_volume_size}"
    }

    tags = {
        Name = "${local.env_tag}-${format("%01d", count.index + 0)}"
    }
}
```

# EC2 instances (main.tf)

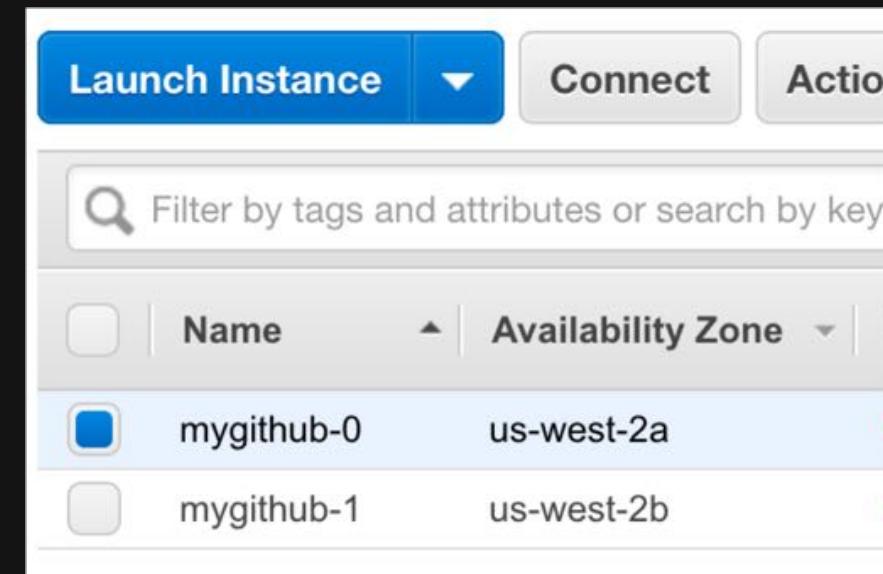
```
resource "aws_instance" "ghe" {
    count          = 2
    ami            = "${data.aws_ami.ghe.id}"
    subnet_id      = "${element(var.subnet_ids, count.index)}"
    instance_type  = "${var.instance_type}"
    vpc_security_group_ids = ["${aws_security_group.ghe.id}"]

    ebs_block_device {
        device_name = "/dev/xvdd"
        volume_type = "gp2"
        volume_size = "${var.data_volume_size}"
    }

    tags = {
        Name = "${local.env_tag}-${format("%01d", count.index + 0)}"
    }
}
```

# EC2 instances Name tag

mygithub.example.com



A screenshot of the AWS EC2 Instances page. At the top, there are buttons for "Launch Instance", "Connect", and "Actions". Below that is a search bar with the placeholder "Filter by tags and attributes or search by key...". The main area shows a table with two rows of instance data:

	Name	Availability Zone
<input checked="" type="checkbox"/>	mygithub-0	us-west-2a
<input type="checkbox"/>	mygithub-1	us-west-2b

# EC2 instances Name tag (main.tf)

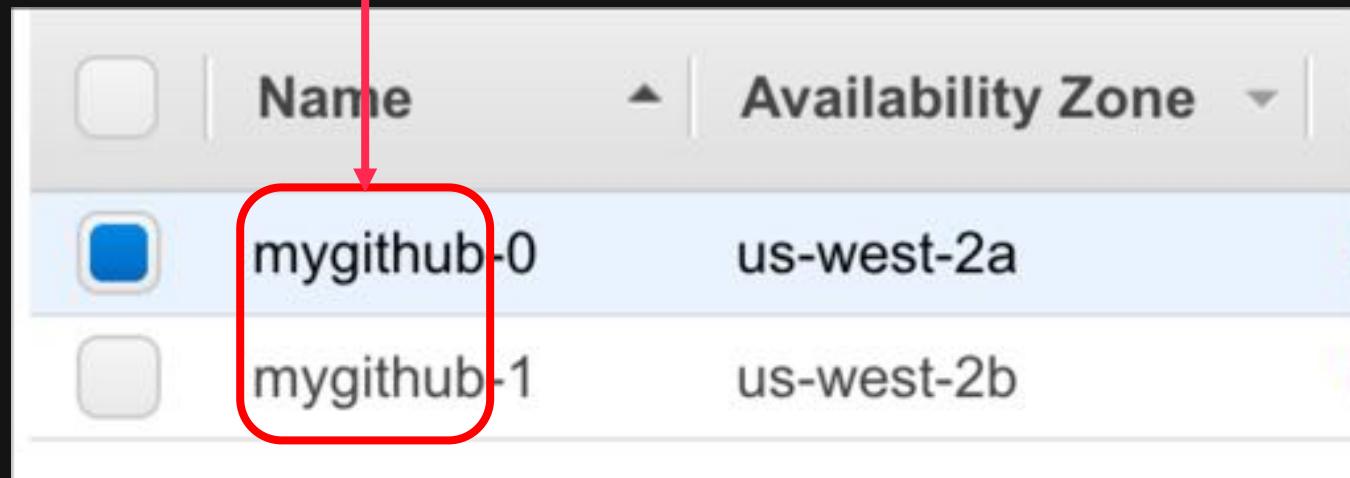
```
locals {  
    env_tag = "${element(split(".", var.ghe_hostname), 0)}"  
}
```

# EC2 instances Name tag (main.tf)

```
locals {  
    env_tag = "${element(split(".", var.ghe_hostname), 0)}"  
}
```

# EC2 instances Name tag (main.tf)

```
tags = {  
  Name = "${local.env_tag}-${format("%01d", count.index + 0)}"  
}
```



	Name	Availability Zone
<input checked="" type="checkbox"/>	mygithub-0	us-west-2a
<input type="checkbox"/>	mygithub-1	us-west-2b

# EC2 instances Name tag (main.tf)

```
tags = {  
  Name = "${local.env_tag}-${format("%01d", count.index + 0)}"  
}
```

	Name	Availability Zone
<input checked="" type="checkbox"/>	mygithub-0	us-west-2a
<input type="checkbox"/>	mygithub-1	us-west-2b

# Security Group and Rules (sg.tf)

```
resource "aws_security_group" "ghe" {
    name      = "ghe-sg"
    vpc_id    = "${var.vpc_id}"
    description = "GitHub Enterprise Appliance Security Group"

    tags = {
        Name      = "${local.env_tag}-appliance SG"
    }
}

resource "aws_security_group_rule" "allow_ssh" {
    type          = "ingress"
    from_port     = 22
    to_port       = 22
    protocol      = "tcp"
    cidr_blocks   = ["0.0.0.0/0"]
    security_group_id = "${aws_security_group.ghe.id}"
}

...
```

# Elastic IP addresses (eip.tf)

```
resource "aws_eip" "ghe_appliance" {
    count      = 2
    instance  = "${element(aws_instance.ghe_appliance.*.id, count.index)}"  
    vpc       = true
}
```

# Route 53 DNS (r53.tf)

```
resource "aws_route53_zone" "ghe" {
  name = "${var.ghe_hostname}"
}

resource "aws_route53_record" "ghe_hostname" {
  zone_id = "${aws_route53_zone.ghe.zone_id}"
  name    = "${var.ghe_hostname}"
  type    = "A"
  ttl     = "300"
  records = ["${var.primary == "0" ? aws_eip.ghe.0.public_ip : aws_eip.ghe.1.public_ip}"]
}

resource "aws_route53_record" "ghe_wildcard" {
  zone_id = "${aws_route53_zone.ghe.zone_id}"
  name    = "*.${var.ghe_hostname}"
  type    = "A"
  ttl     = "300"
  records = ["${var.primary == "0" ? aws_eip.ghe.0.public_ip : aws_eip.ghe.1.public_ip}"]
}
```

# Route 53 DNS (r53.tf)

```
records = ["${var.primary == "0" ? aws_eip.ghe.0.public_ip : aws_eip.ghe.1.public_ip}"]
```

condition

true  
value

false  
value

# Agenda

1. GitHub Enterprise HA on AWS
2. Defining infrastructure with Terraform
3. Using API & Terraform for configuration
4. Future enhancements

Setup GitHub Enterprise

Enterprise

# Install GitHub Enterprise

To finish the installation of GitHub Enterprise, upload your license and provide a password. This password will be used to access the management console as well as the API. SSH administrative access uses authorized SSH keys you've added instead of this password.

 License file  
Your license file must be in zip or tgz.

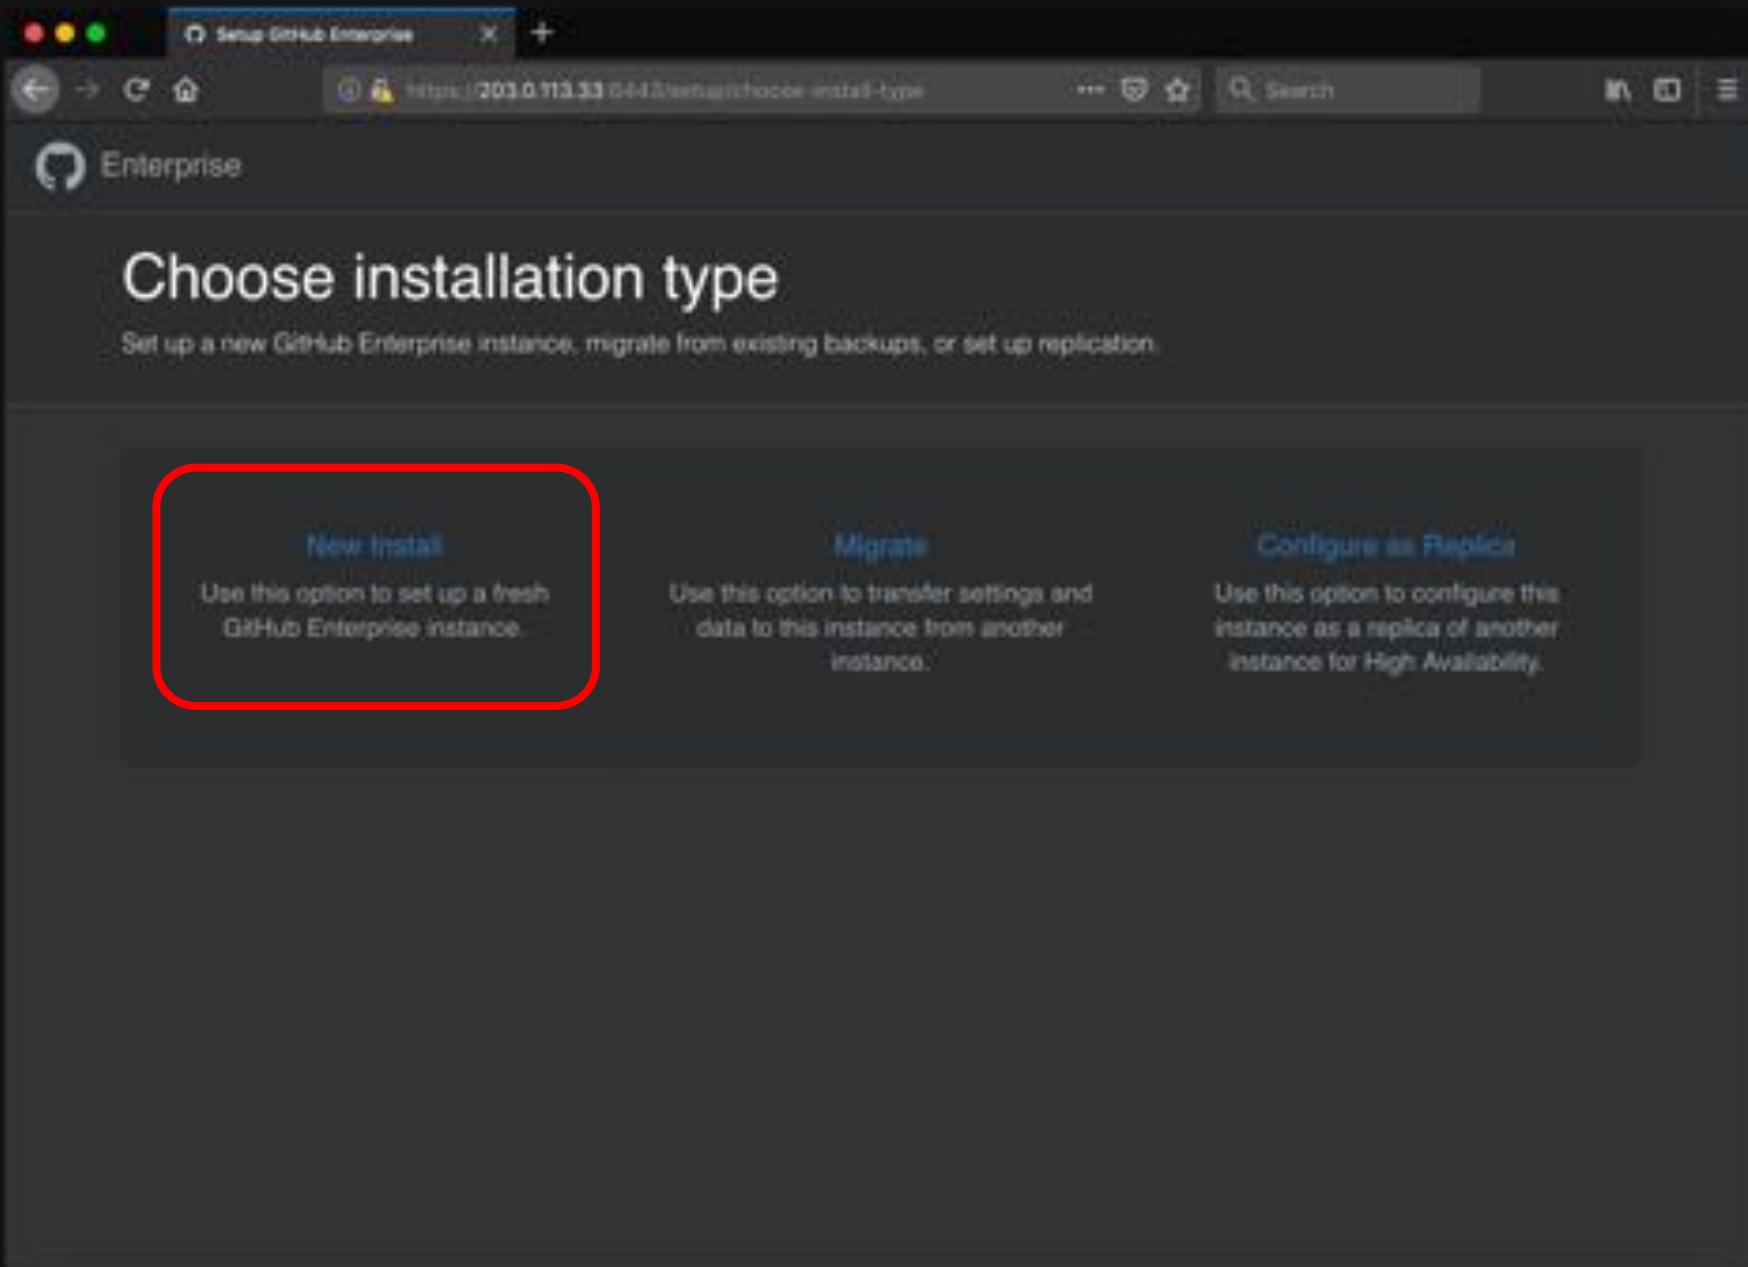
New password

Confirm password

Passwords must be at least 7 characters long and include at least one number and one upper case letter.

Finish installation

 Don't have your license? Download it here.



Setup GitHub Enterprise

https://123.0.113.33:6443/settings

Settings

Hostname

SSH access

Hostname

Time

Authentication

Privacy

Pages

Email

Monitoring

Rate limiting

Hostname

example.com, \*.example.com

mygithub.example.com

Test domain settings

Subdomain isolation (recommended)

Subdomain isolation is an important security feature used to separate user-supplied content from other portions of your GitHub Enterprise installation. A valid domain name, not an IP address, must be set as the Hostname and a wildcard DNS entry must be added for your host.

Time

Save settings

Primary NTP server

ubuntu.pool.ntp.org

Secondary NTP server (optional)

ubuntu.pool.ntp.org

The NTP servers configured here will determine which servers are used to synchronize time with. If the time servers use internal hostnames, you must specify custom DNS nameservers that can resolve them. ⓘ

Authentication

Setup GitHub Enterprise

Enterprise

# Configuring your instance

Your settings are being applied and your instance is restarting at <https://mygithub.example.com>

Preparing storage device	
Updating configuration	
Reloading system services	
Running migrations	
Reloading application services	

[Visit your instance](#)

# Existing appliance settings export (API)

```
curl -sL 'https://api_key:Octocat1@github.example.com:8443/setup/api/settings' \  
 >settings.json
```

# Existing appliance settings export (API)

```
{  
  "enterprise": {  
    "private_mode": true,  
    "public_pages": false,  
    "subdomain_isolation": true,  
    "signup_enabled": false,  
    "github_hostname": "github.example.com",  
    "identicons_host": "dotcom",  
    "http_proxy": null,  
    "http_noproxy": null,  
    "auth_mode": "saml",  
    "builtin_auth_fallback": false,  
    "expire_sessions": false,  
    "admin_password": null,  
    "configuration_id": 1234565432,  
    "configuration_run_count": 5,  
    "avatar": null,  
    ...  
  }  
}
```

# Verify the new appliance is ready

```
curl -skL 'https://203.0.113.33:8443/setup/start' \  
| grep "Setup GitHub Enterprise"
```

# Verify the new appliance is ready

```
curl -skL 'https://203.0.113.33:8443/setup/start' \  
| grep "Setup GitHub Enterprise"
```

# Upload license and set a password (API)

```
curl -skL -X POST 'https://203.0.113.33:8443/setup/api/start' \
-F license=@/tmp/github-enterprise.ghl \
-F "password=reInvent18"
```

# Modify new appliance settings (API)

```
curl -skL -X PUT 'https://api_key:reInvent18@203.0.113.33:8443/setup/api/settings' \
--data-urlencode "settings=`cat /tmp/settings.json`"
```

# Start a configuration process (API)

```
curl -skL -X POST 'https://api_key:reInvent18@203.0.113.33:8443/setup/api/configure'
```

## Parameter details

Name

mygithub/ghe\_password

Description- *Optional*

Type

 String

Any string value.

 StringList

Separate strings using commas.

 SecureStringEncrypt sensitive data using the KMS keys for  
your account.

Create parameter

## Parameter details

Name

mygithub/ghe\_password

Description- *Optional*

Type

 String

Any string value.

 StringList

Separate strings using commas.

 SecureStringEncrypt sensitive data using the KMS keys for  
your account.

# Get Parameter Store password (main.tf)

```
data "aws_ssm_parameter" "ghe_password" {  
    name = "/${local.env_tag}/ghe_password"  
}
```

# Get Parameter Store password (main.tf)

```
locals {  
    env_tag      = "${element(split(".", var.ghe_hostname), 0)}"  
    → ghe_password = "${data.aws_ssm_parameter.ghe_password.value}"  
}
```

# Execute API commands (main.tf)

```
resource "null_resource" "ghe_config" {
    // Check that GHE is ready to take API requests
    provisioner "local-exec" {
        command = "until [ $(curl -sSL 'https://${aws_eip.ghe.0.public_ip}:8443/setup/start') | grep -o 'Setup GitHub...' ]"
    }

    // Upload license and set Management Console password
    provisioner "local-exec" {
        command = "sleep 5 && curl -sSL -X POST 'https://${aws_eip.ghe.0.public_ip}:8443/setup/api/start' -F license..."
    }

    // Upload JSON configuration
    provisioner "local-exec" {
        command = "sleep 10 && curl -sSL -X PUT 'https://api_key:${local.ghe_password}@${aws_eip.ghe.0.public_ip}:8443/api/config' -d @${local.ghe_config}"
    }

    // Start configuring
    provisioner "local-exec" {
        command = "sleep 10 && curl -sSL -X POST 'https://api_key:${local.ghe_password}@${aws_eip.ghe.0.public_ip}:8443/api/config/start'"
    }
}
```

# Module repo contents

```
├── ami.tf  
├── eip.tf  
├── main.tf  
├── r53.tf  
└── sg.tf  
└── variables.tf
```

# "live" repo contents

```
└── github-enterprise.ghl
    ├── main.tf
    └── settings.json
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    instance_type      = "r4.large"
    data_volume_size = "500"
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings  = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    instance_type  = "r4.large"
    data_volume_size = "500"
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings  = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    instance_type  = "r4.large"
    data_volume_size = "500"
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    instance_type      = "r4.large"
    data_volume_size = "500"
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    instance_type = "r4.large"
    data_volume_size = "500"
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings = "settings.json"

    primary = "0"

    [
        vpc_id      = "vpc-123b1234"
        subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

        instance_type      = "r4.large"
        data_volume_size = "500"
    ]
}
```

```
provider "aws" {
    region = "us-west-2"
}

module "ghe" {
    source = "git::ssh://git@github.com/felipe-ramirez/reinvent18?ref=v0.1.0"

    ghe_version  = "2.15.1"
    ghe_hostname = "mygithub.example.com"
    ghe_license   = "github-enterprise.ghl"
    ghe_settings  = "settings.json"

    primary = "0"

    vpc_id      = "vpc-123b1234"
    subnet_ids  = ["subnet-546c8a45", "subnet-23c45d6e"]

    [
        instance_type      = "r4.large"
        data_volume_size = "500"
    ]
}
```

# Agenda

1. GitHub Enterprise HA on AWS
2. Defining infrastructure with Terraform
3. Using API & Terraform for configuration
4. Future enhancements

Automate HA  
configuration on replica  
appliance

Use AWS CodePipeline  
for applying updates

Automated daily testing  
of backup data



[github.com/felipe-ramirez/reinvent18](https://github.com/felipe-ramirez/reinvent18)

# Thank you!

**Bryan Cross**

bryancross@github.com

**Felipe Ramirez**

felipe.ramirez@coxautoinc.com